

```

from torchvision import datasets, transforms
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.nn.init as init
import numpy as np
from torch.utils.data import DataLoader
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from torch.utils.data import TensorDataset, DataLoader
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import pandas as pd
import time
from torch.utils.data import random_split, DataLoader
import torchvision.transforms as transforms
import torchvision.datasets as datasets

```

▼ Logistic Regression

```

transform = transforms.Compose([
    transforms.ToTensor(), # convert image to PyTorch tensor
    transforms.Normalize((0.0,), (1.0,)) # normalize to [0,1]
])

mnist_train = datasets.MNIST(
    root='data', train=True, transform=transform, download=True)
mnist_test = datasets.MNIST(
    root='data', train=False, transform=transform, download=True)

```

```

X = []
y = []

for img, label in mnist_train:
    if label in [0, 1]:
        X.append(img.view(-1)) # flatten 28x28 -> 784
        y.append(label)

X = torch.stack(X)
y = torch.tensor(y)
X_train, X_temp, y_train, y_temp = train_test_split(
    X, y, test_size=0.4, stratify=y)
X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, stratify=y_temp)
print(
    f"Train size: {X_train.shape[0]}, Val size: {X_val.shape[0]}, Test size: {X_test.shape[0]}")

Train size: 7599, Val size: 2533, Test size: 2533

```

```

train_loader = DataLoader(TensorDataset(
    X_train, y_train), batch_size=64, shuffle=True)
val_loader = DataLoader(TensorDataset(X_val, y_val), batch_size=64)
test_loader = DataLoader(TensorDataset(X_test, y_test), batch_size=64)

```

```

def sigmoid(z):
    return 1 / (1 + torch.exp(-z))

def binary_cross_entropy(y_hat, y):
    eps = 1e-15
    y_hat = torch.clamp(y_hat, eps, 1 - eps) # prevent log(0)
    # y_hat < eps = eps
    # y_hat > 1 - eps = 1 - eps
    return -torch.mean(y * torch.log(y_hat) + (1 - y) * torch.log(1 - y_hat))

# Number of features = 784 pixels per image
n_features = 784

# Initialize weights and bias

```

```
W = torch.zeros((n_features, 1), requires_grad=True)
b = torch.zeros(1, requires_grad=True)
```

```
# --- Hyperparameters and trackers ---
learning_rate = 0.01
epochs = 100
train_losses = []
val_losses = []
val_accuracies = []
train_accuracies = []

# --- Training loop ---
start_time = time.time()
for epoch in range(epochs):
    epoch_loss = 0.0
    correct_train = 0
    total_train = 0

    # ----- TRAINING PHASE -----
    for X_batch, y_batch in train_loader:
        # Forward pass
        y_pred = sigmoid(X_batch @ W + b)
        loss = binary_cross_entropy(y_pred, y_batch.unsqueeze(1).float())

        # Backward pass
        loss.backward()

        # Gradient update (manual SGD)
        with torch.no_grad():
            W -= learning_rate * W.grad
            b -= learning_rate * b.grad

        # Reset gradients
        W.grad.zero_()
        b.grad.zero_()

        # Accumulate batch loss
        epoch_loss += loss.item()
        # Compute training accuracy per batch
        preds = (y_pred > 0.5).int().squeeze()
        correct_train += (preds == y_batch).sum().item()
        total_train += y_batch.size(0)

    # Compute average training loss for this epoch
    avg_train_loss = epoch_loss / len(train_loader)
    train_losses.append(avg_train_loss)
    train_accuracies.append(correct_train / total_train)

    # ----- VALIDATION PHASE -----
    with torch.no_grad():
        val_loss_total = 0.0
        correct = 0
        total = 0

        for X_val_batch, y_val_batch in val_loader:
            y_val_pred = sigmoid(X_val_batch @ W + b)
            val_loss_total += binary_cross_entropy(
                y_val_pred, y_val_batch.unsqueeze(1).float()).item()

            # Compute accuracy
            y_val_pred_label = (y_val_pred > 0.5).int().squeeze()
            correct += (y_val_pred_label == y_val_batch).sum().item()
            total += y_val_batch.size(0)

        avg_val_loss = val_loss_total / len(val_loader)
        val_acc = correct / total

        val_losses.append(avg_val_loss)
        val_accuracies.append(val_acc)

    # Print progress for this epoch
    print(f"Epoch {epoch+1}/{epochs} | "
          f"Train Loss: {avg_train_loss:.4f} | "
          f"Val Loss: {avg_val_loss:.4f} | "
          f"Train Acc: {train_accuracies[-1]:.4f} | "
          f"Val Acc: {val_accuracies[-1]:.4f}")
```

```

end_time = time.time()
lr_run_time = end_time - start_time

```

Epoch 1/100	Train Loss: 0.2320	Val Loss: 0.1053	Train Acc: 0.9904	Val Acc: 0.9972
Epoch 2/100	Train Loss: 0.0795	Val Loss: 0.0618	Train Acc: 0.9962	Val Acc: 0.9976
Epoch 3/100	Train Loss: 0.0532	Val Loss: 0.0455	Train Acc: 0.9963	Val Acc: 0.9976
Epoch 4/100	Train Loss: 0.0415	Val Loss: 0.0369	Train Acc: 0.9968	Val Acc: 0.9972
Epoch 5/100	Train Loss: 0.0348	Val Loss: 0.0315	Train Acc: 0.9970	Val Acc: 0.9972
Epoch 6/100	Train Loss: 0.0303	Val Loss: 0.0277	Train Acc: 0.9972	Val Acc: 0.9972
Epoch 7/100	Train Loss: 0.0272	Val Loss: 0.0249	Train Acc: 0.9972	Val Acc: 0.9976
Epoch 8/100	Train Loss: 0.0248	Val Loss: 0.0228	Train Acc: 0.9972	Val Acc: 0.9976
Epoch 9/100	Train Loss: 0.0228	Val Loss: 0.0211	Train Acc: 0.9971	Val Acc: 0.9976
Epoch 10/100	Train Loss: 0.0213	Val Loss: 0.0197	Train Acc: 0.9972	Val Acc: 0.9976
Epoch 11/100	Train Loss: 0.0200	Val Loss: 0.0185	Train Acc: 0.9972	Val Acc: 0.9976
Epoch 12/100	Train Loss: 0.0190	Val Loss: 0.0175	Train Acc: 0.9972	Val Acc: 0.9976
Epoch 13/100	Train Loss: 0.0181	Val Loss: 0.0167	Train Acc: 0.9972	Val Acc: 0.9976
Epoch 14/100	Train Loss: 0.0172	Val Loss: 0.0159	Train Acc: 0.9972	Val Acc: 0.9976
Epoch 15/100	Train Loss: 0.0165	Val Loss: 0.0153	Train Acc: 0.9972	Val Acc: 0.9976
Epoch 16/100	Train Loss: 0.0159	Val Loss: 0.0147	Train Acc: 0.9974	Val Acc: 0.9976
Epoch 17/100	Train Loss: 0.0154	Val Loss: 0.0142	Train Acc: 0.9974	Val Acc: 0.9976
Epoch 18/100	Train Loss: 0.0149	Val Loss: 0.0137	Train Acc: 0.9974	Val Acc: 0.9976
Epoch 19/100	Train Loss: 0.0144	Val Loss: 0.0133	Train Acc: 0.9974	Val Acc: 0.9976
Epoch 20/100	Train Loss: 0.0140	Val Loss: 0.0129	Train Acc: 0.9974	Val Acc: 0.9976
Epoch 21/100	Train Loss: 0.0137	Val Loss: 0.0126	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 22/100	Train Loss: 0.0133	Val Loss: 0.0122	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 23/100	Train Loss: 0.0129	Val Loss: 0.0119	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 24/100	Train Loss: 0.0127	Val Loss: 0.0117	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 25/100	Train Loss: 0.0124	Val Loss: 0.0114	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 26/100	Train Loss: 0.0122	Val Loss: 0.0112	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 27/100	Train Loss: 0.0119	Val Loss: 0.0109	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 28/100	Train Loss: 0.0116	Val Loss: 0.0107	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 29/100	Train Loss: 0.0114	Val Loss: 0.0105	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 30/100	Train Loss: 0.0113	Val Loss: 0.0103	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 31/100	Train Loss: 0.0111	Val Loss: 0.0102	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 32/100	Train Loss: 0.0109	Val Loss: 0.0100	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 33/100	Train Loss: 0.0107	Val Loss: 0.0098	Train Acc: 0.9975	Val Acc: 0.9976
Epoch 34/100	Train Loss: 0.0106	Val Loss: 0.0097	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 35/100	Train Loss: 0.0104	Val Loss: 0.0096	Train Acc: 0.9976	Val Acc: 0.9976
Epoch 36/100	Train Loss: 0.0102	Val Loss: 0.0094	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 37/100	Train Loss: 0.0101	Val Loss: 0.0093	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 38/100	Train Loss: 0.0100	Val Loss: 0.0092	Train Acc: 0.9976	Val Acc: 0.9976
Epoch 39/100	Train Loss: 0.0099	Val Loss: 0.0091	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 40/100	Train Loss: 0.0097	Val Loss: 0.0090	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 41/100	Train Loss: 0.0096	Val Loss: 0.0089	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 42/100	Train Loss: 0.0095	Val Loss: 0.0088	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 43/100	Train Loss: 0.0094	Val Loss: 0.0087	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 44/100	Train Loss: 0.0093	Val Loss: 0.0086	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 45/100	Train Loss: 0.0092	Val Loss: 0.0085	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 46/100	Train Loss: 0.0091	Val Loss: 0.0084	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 47/100	Train Loss: 0.0091	Val Loss: 0.0083	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 48/100	Train Loss: 0.0089	Val Loss: 0.0082	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 49/100	Train Loss: 0.0088	Val Loss: 0.0082	Train Acc: 0.9979	Val Acc: 0.9976
Epoch 50/100	Train Loss: 0.0087	Val Loss: 0.0081	Train Acc: 0.9978	Val Acc: 0.9976
Epoch 51/100	Train Loss: 0.0086	Val Loss: 0.0080	Train Acc: 0.9979	Val Acc: 0.9976
Epoch 52/100	Train Loss: 0.0085	Val Loss: 0.0079	Train Acc: 0.9980	Val Acc: 0.9976
Epoch 53/100	Train Loss: 0.0085	Val Loss: 0.0079	Train Acc: 0.9980	Val Acc: 0.9976
Epoch 54/100	Train Loss: 0.0084	Val Loss: 0.0078	Train Acc: 0.9982	Val Acc: 0.9976
Epoch 55/100	Train Loss: 0.0083	Val Loss: 0.0077	Train Acc: 0.9980	Val Acc: 0.9976
Epoch 56/100	Train Loss: 0.0083	Val Loss: 0.0077	Train Acc: 0.9980	Val Acc: 0.9976
Epoch 57/100	Train Loss: 0.0082	Val Loss: 0.0076	Train Acc: 0.9984	Val Acc: 0.9976
Epoch 58/100	Train Loss: 0.0082	Val Loss: 0.0076	Train Acc: 0.9982	Val Acc: 0.9976

```

plt.figure(figsize=(8, 5))
plt.plot(train_losses, label='Training Loss')
plt.plot(val_losses, label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training vs Validation Loss')
plt.legend()
plt.show()

plt.figure(figsize=(8, 5))
plt.plot(train_accuracies, label='Training Accuracy')
plt.plot(val_accuracies, label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training vs Validation Accuracy')
plt.legend()
plt.show()

with torch.no_grad():


```

```
y_test_pred = sigmoid(X_test @ W + b)
y_test_pred_label = (y_test_pred > 0.5).int().squeeze()


test_acc = (y_test_pred_label == y_test).float().mean()
print(f"Test Accuracy: {test_acc:.4f}")

cm = confusion_matrix(y_test, y_test_pred_label)
ConfusionMatrixDisplay(cm).plot()
plt.show()
```


Training vs Validation Loss



Training vs Validation Accuracy



Test Accuracy: 0.9992



```
model_sklearn = LogisticRegression(max_iter=100, solver='liblinear')
model_sklearn.fit(X_train, y_train)
```

```

train_acc_sklearn = model_sklearn.score(X_train, y_train)
val_acc_sklearn = model_sklearn.score(X_val, y_val)
test_acc_sklearn = model_sklearn.score(X_test, y_test)

print(
    f"Sklearn Logistic Regression - Training Accuracy: {train_acc_sklearn:.4f}")
print(
    f"Sklearn Logistic Regression - Validation Accuracy: {val_acc_sklearn:.4f}")
print(f"Sklearn Logistic Regression - Test Accuracy: {test_acc_sklearn:.4f}")

Sklearn Logistic Regression - Training Accuracy: 1.0000
Sklearn Logistic Regression - Validation Accuracy: 0.9980
Sklearn Logistic Regression - Test Accuracy: 0.9992

```

```



manual_final_train_acc = train_accuracies[-1]
manual_final_val_acc = val_accuracies[-1]

plt.figure(figsize=(8, 5))
plt.plot(range(1, len(train_accuracies)+1), train_accuracies,
         label='Manual Training Accuracy', linewidth=2)
plt.axhline(y=train_acc_sklearn, color='orange',
            linestyle='--', label='Sklearn Training Accuracy')

plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training Accuracy: Manual vs Sklearn Logistic Regression')
plt.legend()
plt.grid(True)
plt.show()

plt.figure(figsize=(8, 5))
plt.plot(range(1, len(val_accuracies)+1), val_accuracies,
         label='Manual Validation Accuracy', linewidth=2)
plt.axhline(y=val_acc_sklearn, color='orange', linestyle='--',
            label='Sklearn Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Validation Accuracy: Manual vs Sklearn Logistic Regression')
plt.legend()
plt.grid(True)
plt.show()

```




```

labels = ['Train Accuracy', 'Validation Accuracy', 'Test Accuracy']
manual_values = [
    manual_final_train_acc,
    manual_final_val_acc,
    (y_test_pred_label == y_test).float().mean().item()
]
sklearn_values = [
    train_acc_sklearn,
    val_acc_sklearn,
    model_sklearn.score(X_test, y_test)
]
x = np.arange(len(labels))
width = 0.35

plt.figure(figsize=(8, 5))
plt.bar(x - width/2, manual_values, width, label='Manual')
plt.bar(x + width/2, sklearn_values, width, label='Sklearn')
plt.ylabel('Accuracy')
plt.title('Manual vs Sklearn Logistic Regression Performance')
plt.xticks(x, labels)
plt.ylim(0.95, 1.01)
plt.legend()
plt.grid(axis='y')
plt.show()

```

Manual vs Sklearn Logistic Regression Performance



```

transform = transforms.Compose([transforms.ToTensor()])
mnist_train = datasets.MNIST(
    root='data', train=True, transform=transform, download=True)
# Filter digits 0 and 1
X = []
y = []
for img, label in mnist_train:
    if label in [0, 1]:
        X.append(img.view(-1))
        y.append(label)

X = torch.stack(X)
y = torch.tensor(y, dtype=torch.float32)
X_train, X_temp, y_train, y_temp = train_test_split(
    X, y, test_size=0.4, stratify=y)
X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, stratify=y_temp)
print(f"Train set size: {X_train.shape[0]}")
print(f"Validation set size: {X_val.shape[0]}")
print(f"Test set size: {X_test.shape[0]}")
train_loader = DataLoader(TensorDataset(
    X_train, y_train), batch_size=64, shuffle=True)
val_loader = DataLoader(TensorDataset(X_val, y_val), batch_size=64)
test_loader = DataLoader(TensorDataset(X_test, y_test), batch_size=64)
model = nn.Linear(784, 1)
criterion = nn.BCEWithLogitsLoss()      # combines sigmoid + BCE for stability
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)

Train set size: 7599
Validation set size: 2533
Test set size: 2533

```

```

epochs = 100
train_losses, val_losses = [], []
train_accuracies, val_accuracies = [], []

for epoch in range(epochs):
    model.train()
    running_loss = 0.0
    correct_train = 0
    total_train = 0

    # TRAIN
    for X_batch, y_batch in train_loader:
        optimizer.zero_grad()
        outputs = model(X_batch).squeeze()
        loss = criterion(outputs, y_batch)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

```

```

preds = (torch.sigmoid(outputs) > 0.5).int()
correct_train += (preds == y_batch.int()).sum().item()
total_train += y_batch.size(0)

avg_train_loss = running_loss / len(train_loader)
train_acc = correct_train / total_train
train_losses.append(avg_train_loss)
train_accuracies.append(train_acc)

# VALIDATION
model.eval()
val_loss = 0.0
correct_val = 0
total_val = 0
with torch.no_grad():
    for X_val_batch, y_val_batch in val_loader:
        outputs_val = model(X_val_batch).squeeze()
        loss = criterion(outputs_val, y_val_batch)
        val_loss += loss.item()

        preds_val = (torch.sigmoid(outputs_val) > 0.5).int()
        correct_val += (preds_val == y_val_batch.int()).sum().item()
        total_val += y_val_batch.size(0)

    avg_val_loss = val_loss / len(val_loader)
    val_acc = correct_val / total_val
    val_losses.append(avg_val_loss)
    val_accuracies.append(val_acc)

print(f"Epoch {epoch+1}/{epochs} | "
      f"Train Loss: {avg_train_loss:.4f} | Val Loss: {avg_val_loss:.4f} | "
      f"Train Acc: {train_acc:.4f} | Val Acc: {val_acc:.4f}")

```

Epoch 1/100	Train Loss: 0.2408	Val Loss: 0.1076	Train Acc: 0.9745	Val Acc: 0.9945
Epoch 2/100	Train Loss: 0.0796	Val Loss: 0.0640	Train Acc: 0.9968	Val Acc: 0.9945
Epoch 3/100	Train Loss: 0.0526	Val Loss: 0.0480	Train Acc: 0.9968	Val Acc: 0.9953
Epoch 4/100	Train Loss: 0.0406	Val Loss: 0.0395	Train Acc: 0.9972	Val Acc: 0.9953
Epoch 5/100	Train Loss: 0.0338	Val Loss: 0.0342	Train Acc: 0.9974	Val Acc: 0.9957
Epoch 6/100	Train Loss: 0.0292	Val Loss: 0.0305	Train Acc: 0.9975	Val Acc: 0.9957
Epoch 7/100	Train Loss: 0.0260	Val Loss: 0.0279	Train Acc: 0.9976	Val Acc: 0.9957
Epoch 8/100	Train Loss: 0.0235	Val Loss: 0.0258	Train Acc: 0.9978	Val Acc: 0.9957
Epoch 9/100	Train Loss: 0.0216	Val Loss: 0.0242	Train Acc: 0.9979	Val Acc: 0.9957
Epoch 10/100	Train Loss: 0.0201	Val Loss: 0.0228	Train Acc: 0.9979	Val Acc: 0.9957
Epoch 11/100	Train Loss: 0.0188	Val Loss: 0.0217	Train Acc: 0.9980	Val Acc: 0.9961
Epoch 12/100	Train Loss: 0.0177	Val Loss: 0.0208	Train Acc: 0.9980	Val Acc: 0.9961
Epoch 13/100	Train Loss: 0.0168	Val Loss: 0.0199	Train Acc: 0.9980	Val Acc: 0.9961
Epoch 14/100	Train Loss: 0.0160	Val Loss: 0.0192	Train Acc: 0.9980	Val Acc: 0.9961
Epoch 15/100	Train Loss: 0.0153	Val Loss: 0.0186	Train Acc: 0.9980	Val Acc: 0.9961
Epoch 16/100	Train Loss: 0.0146	Val Loss: 0.0181	Train Acc: 0.9980	Val Acc: 0.9961
Epoch 17/100	Train Loss: 0.0141	Val Loss: 0.0176	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 18/100	Train Loss: 0.0136	Val Loss: 0.0171	Train Acc: 0.9980	Val Acc: 0.9961
Epoch 19/100	Train Loss: 0.0132	Val Loss: 0.0167	Train Acc: 0.9980	Val Acc: 0.9961
Epoch 20/100	Train Loss: 0.0127	Val Loss: 0.0163	Train Acc: 0.9980	Val Acc: 0.9961
Epoch 21/100	Train Loss: 0.0124	Val Loss: 0.0160	Train Acc: 0.9980	Val Acc: 0.9961
Epoch 22/100	Train Loss: 0.0120	Val Loss: 0.0157	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 23/100	Train Loss: 0.0117	Val Loss: 0.0154	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 24/100	Train Loss: 0.0114	Val Loss: 0.0152	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 25/100	Train Loss: 0.0111	Val Loss: 0.0149	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 26/100	Train Loss: 0.0109	Val Loss: 0.0147	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 27/100	Train Loss: 0.0106	Val Loss: 0.0145	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 28/100	Train Loss: 0.0104	Val Loss: 0.0143	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 29/100	Train Loss: 0.0102	Val Loss: 0.0141	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 30/100	Train Loss: 0.0100	Val Loss: 0.0139	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 31/100	Train Loss: 0.0098	Val Loss: 0.0137	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 32/100	Train Loss: 0.0097	Val Loss: 0.0136	Train Acc: 0.9982	Val Acc: 0.9961
Epoch 33/100	Train Loss: 0.0095	Val Loss: 0.0134	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 34/100	Train Loss: 0.0093	Val Loss: 0.0133	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 35/100	Train Loss: 0.0092	Val Loss: 0.0131	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 36/100	Train Loss: 0.0090	Val Loss: 0.0130	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 37/100	Train Loss: 0.0089	Val Loss: 0.0129	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 38/100	Train Loss: 0.0088	Val Loss: 0.0128	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 39/100	Train Loss: 0.0086	Val Loss: 0.0127	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 40/100	Train Loss: 0.0085	Val Loss: 0.0126	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 41/100	Train Loss: 0.0084	Val Loss: 0.0124	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 42/100	Train Loss: 0.0083	Val Loss: 0.0123	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 43/100	Train Loss: 0.0082	Val Loss: 0.0123	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 44/100	Train Loss: 0.0081	Val Loss: 0.0122	Train Acc: 0.9984	Val Acc: 0.9961
Epoch 45/100	Train Loss: 0.0080	Val Loss: 0.0121	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 46/100	Train Loss: 0.0079	Val Loss: 0.0120	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 47/100	Train Loss: 0.0078	Val Loss: 0.0119	Train Acc: 0.9983	Val Acc: 0.9961
Epoch 48/100	Train Loss: 0.0077	Val Loss: 0.0118	Train Acc: 0.9983	Val Acc: 0.9961


Epoch 49/100	Train Loss: 0.0077	Val Loss: 0.0118	Train Acc: 0.9983	Val Acc: 0.9964
Epoch 50/100	Train Loss: 0.0076	Val Loss: 0.0117	Train Acc: 0.9983	Val Acc: 0.9964
Epoch 51/100	Train Loss: 0.0076	Val Loss: 0.0116	Train Acc: 0.9983	Val Acc: 0.9964
Epoch 52/100	Train Loss: 0.0074	Val Loss: 0.0115	Train Acc: 0.9983	Val Acc: 0.9964
Epoch 53/100	Train Loss: 0.0073	Val Loss: 0.0115	Train Acc: 0.9983	Val Acc: 0.9964
Epoch 54/100	Train Loss: 0.0073	Val Loss: 0.0114	Train Acc: 0.9984	Val Acc: 0.9964
Epoch 55/100	Train Loss: 0.0072	Val Loss: 0.0114	Train Acc: 0.9984	Val Acc: 0.9964
Epoch 56/100	Train Loss: 0.0071	Val Loss: 0.0113	Train Acc: 0.9984	Val Acc: 0.9964
Epoch 57/100	Train Loss: 0.0071	Val Loss: 0.0112	Train Acc: 0.9986	Val Acc: 0.9964
Epoch 58/100	Train Loss: 0.0070	Val Loss: 0.0112	Train Acc: 0.9986	Val Acc: 0.9964

```

plt.figure(figsize=(8, 5))
plt.plot(train_losses, label='Training Loss')
plt.plot(val_losses, label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('PyTorch Logistic Regression - Training vs Validation Loss')
plt.legend()
plt.grid(alpha=0.3)
plt.show()

best_logistic_accuracy = max(val_accuracies)

```



▼ Soft Max Regression

```

transform = transforms.Compose([
    transforms.ToTensor() # Convert to tensor and scale to [0,1]
])
mnist_train = datasets.MNIST(
    root='data', train=True, transform=transform, download=True)
mnist_test = datasets.MNIST(
    root='data', train=False, transform=transform, download=True)
X = torch.stack([img.view(-1) for img, _ in mnist_train])
y = torch.tensor([label for _, label in mnist_train])
X_train, X_temp, y_train, y_temp = train_test_split(
    X, y, test_size=0.4, stratify=y)
X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, stratify=y_temp)
print(f"Train set size: {X_train.shape[0]}")
print(f"Validation set size: {X_val.shape[0]}")
print(f"Test set size: {X_test.shape[0]}")
batch_size = 64
train_loader = DataLoader(TensorDataset(
    X_train, y_train), batch_size=batch_size, shuffle=True)
val_loader = DataLoader(TensorDataset(X_val, y_val), batch_size=batch_size)
test_loader = DataLoader(TensorDataset(X_test, y_test), batch_size=batch_size)

```

```

Train set size: 36000
Validation set size: 12000
Test set size: 12000

```

```

def softmax(z):
    exp_z = torch.exp(z - torch.max(z, dim=1, keepdim=True).values)
    return exp_z / exp_z.sum(dim=1, keepdim=True)

def cross_entropy_loss(y_hat, y_true):
    eps = 1e-15
    y_hat = torch.clamp(y_hat, eps, 1 - eps)
    one_hot = torch.nn.functional.one_hot(y_true, num_classes=10).float()
    return -torch.mean(torch.sum(one_hot * torch.log(y_hat), dim=1))

n_features = 784
n_classes = 10
W = torch.zeros((n_features, n_classes), requires_grad=True)
b = torch.zeros(n_classes, requires_grad=True)
learning_rate = 0.01
epochs = 100
train_losses = []
val_losses = []
train_accuracies = []
val_accuracies = []

```

```

# --- Training loop ---
start_time = time.time()
for epoch in range(epochs):
    epoch_loss = 0.0
    correct_train = 0
    total_train = 0

    # ----- TRAINING PHASE -----
    for X_batch, y_batch in train_loader:
        # Forward pass
        y_pred = softmax(X_batch @ W + b)
        loss = cross_entropy_loss(y_pred, y_batch)

        # Backward pass
        loss.backward()

        # Gradient update (manual SGD)
        with torch.no_grad():
            W -= learning_rate * W.grad
            b -= learning_rate * b.grad

        # Reset gradients
        W.grad.zero_()
        b.grad.zero_()

        # Accumulate batch loss
        epoch_loss += loss.item()
        # Compute training accuracy per batch
        preds = y_pred.argmax(dim=1)
        correct_train += (preds == y_batch).sum().item()
        total_train += y_batch.size(0)

    # Compute average training loss for this epoch
    avg_train_loss = epoch_loss / len(train_loader)
    train_acc = correct_train / total_train
    train_losses.append(avg_train_loss)
    train_accuracies.append(train_acc)

    # ----- VALIDATION PHASE -----
    with torch.no_grad():
        val_loss_total = 0.0
        correct = 0
        total = 0

        for X_val_batch, y_val_batch in val_loader:
            y_val_pred = softmax(X_val_batch @ W + b)
            val_loss_total += cross_entropy_loss(y_val_pred,
                                                y_val_batch).item()

            # Compute accuracy
            y_val_pred_label = y_val_pred.argmax(dim=1)
            correct += (y_val_pred_label == y_val_batch).sum().item()

```

```

        total += y_val_batch.size(0)

    avg_val_loss = val_loss_total / len(val_loader)
    val_acc = correct / total

    val_losses.append(avg_val_loss)
    val_accuracies.append(val_acc)

    # Print progress for this epoch
    print(f"Epoch {epoch+1}/{epochs} | "
          f"Train Loss: {avg_train_loss:.4f} | "
          f"Val Loss: {avg_val_loss:.4f} | "
          f"Train Acc: {train_accuracies[-1]:.4f} | "
          f"Val Acc: {val_accuracies[-1]:.4f}")

end_time = time.time()
sm_run_time = end_time - start_time

```

Epoch 1/100	Train Loss: 1.1717	Val Loss: 0.7680	Train Acc: 0.7889	Val Acc: 0.8396
Epoch 2/100	Train Loss: 0.6561	Val Loss: 0.5932	Train Acc: 0.8547	Val Acc: 0.8612
Epoch 3/100	Train Loss: 0.5455	Val Loss: 0.5229	Train Acc: 0.8680	Val Acc: 0.8700
Epoch 4/100	Train Loss: 0.4921	Val Loss: 0.4828	Train Acc: 0.8770	Val Acc: 0.8784
Epoch 5/100	Train Loss: 0.4590	Val Loss: 0.4566	Train Acc: 0.8824	Val Acc: 0.8832
Epoch 6/100	Train Loss: 0.4364	Val Loss: 0.4382	Train Acc: 0.8858	Val Acc: 0.8871
Epoch 7/100	Train Loss: 0.4194	Val Loss: 0.4231	Train Acc: 0.8889	Val Acc: 0.8894
Epoch 8/100	Train Loss: 0.4059	Val Loss: 0.4121	Train Acc: 0.8909	Val Acc: 0.8911
Epoch 9/100	Train Loss: 0.3954	Val Loss: 0.4026	Train Acc: 0.8930	Val Acc: 0.8926
Epoch 10/100	Train Loss: 0.3863	Val Loss: 0.3947	Train Acc: 0.8943	Val Acc: 0.8945
Epoch 11/100	Train Loss: 0.3789	Val Loss: 0.3883	Train Acc: 0.8960	Val Acc: 0.8956
Epoch 12/100	Train Loss: 0.3722	Val Loss: 0.3820	Train Acc: 0.8972	Val Acc: 0.8969
Epoch 13/100	Train Loss: 0.3665	Val Loss: 0.3771	Train Acc: 0.8988	Val Acc: 0.8986
Epoch 14/100	Train Loss: 0.3614	Val Loss: 0.3726	Train Acc: 0.9004	Val Acc: 0.8994
Epoch 15/100	Train Loss: 0.3568	Val Loss: 0.3691	Train Acc: 0.9010	Val Acc: 0.8998
Epoch 16/100	Train Loss: 0.3526	Val Loss: 0.3647	Train Acc: 0.9024	Val Acc: 0.9007
Epoch 17/100	Train Loss: 0.3487	Val Loss: 0.3613	Train Acc: 0.9031	Val Acc: 0.9012
Epoch 18/100	Train Loss: 0.3453	Val Loss: 0.3583	Train Acc: 0.9035	Val Acc: 0.9025
Epoch 19/100	Train Loss: 0.3424	Val Loss: 0.3556	Train Acc: 0.9044	Val Acc: 0.9033
Epoch 20/100	Train Loss: 0.3394	Val Loss: 0.3532	Train Acc: 0.9058	Val Acc: 0.9038
Epoch 21/100	Train Loss: 0.3367	Val Loss: 0.3508	Train Acc: 0.9061	Val Acc: 0.9037
Epoch 22/100	Train Loss: 0.3341	Val Loss: 0.3485	Train Acc: 0.9067	Val Acc: 0.9040
Epoch 23/100	Train Loss: 0.3318	Val Loss: 0.3464	Train Acc: 0.9074	Val Acc: 0.9054
Epoch 24/100	Train Loss: 0.3295	Val Loss: 0.3449	Train Acc: 0.9078	Val Acc: 0.9058
Epoch 25/100	Train Loss: 0.3273	Val Loss: 0.3427	Train Acc: 0.9084	Val Acc: 0.9053
Epoch 26/100	Train Loss: 0.3254	Val Loss: 0.3410	Train Acc: 0.9087	Val Acc: 0.9060
Epoch 27/100	Train Loss: 0.3239	Val Loss: 0.3394	Train Acc: 0.9094	Val Acc: 0.9067
Epoch 28/100	Train Loss: 0.3223	Val Loss: 0.3382	Train Acc: 0.9101	Val Acc: 0.9073
Epoch 29/100	Train Loss: 0.3202	Val Loss: 0.3365	Train Acc: 0.9099	Val Acc: 0.9075
Epoch 30/100	Train Loss: 0.3187	Val Loss: 0.3354	Train Acc: 0.9108	Val Acc: 0.9083
Epoch 31/100	Train Loss: 0.3171	Val Loss: 0.3340	Train Acc: 0.9110	Val Acc: 0.9076
Epoch 32/100	Train Loss: 0.3160	Val Loss: 0.3329	Train Acc: 0.9113	Val Acc: 0.9089
Epoch 33/100	Train Loss: 0.3145	Val Loss: 0.3317	Train Acc: 0.9115	Val Acc: 0.9090
Epoch 34/100	Train Loss: 0.3135	Val Loss: 0.3306	Train Acc: 0.9126	Val Acc: 0.9089
Epoch 35/100	Train Loss: 0.3118	Val Loss: 0.3293	Train Acc: 0.9126	Val Acc: 0.9092
Epoch 36/100	Train Loss: 0.3105	Val Loss: 0.3287	Train Acc: 0.9130	Val Acc: 0.9093
Epoch 37/100	Train Loss: 0.3094	Val Loss: 0.3278	Train Acc: 0.9136	Val Acc: 0.9093
Epoch 38/100	Train Loss: 0.3083	Val Loss: 0.3268	Train Acc: 0.9139	Val Acc: 0.9097
Epoch 39/100	Train Loss: 0.3073	Val Loss: 0.3255	Train Acc: 0.9137	Val Acc: 0.9111
Epoch 40/100	Train Loss: 0.3062	Val Loss: 0.3254	Train Acc: 0.9143	Val Acc: 0.9099
Epoch 41/100	Train Loss: 0.3052	Val Loss: 0.3242	Train Acc: 0.9148	Val Acc: 0.9103
Epoch 42/100	Train Loss: 0.3043	Val Loss: 0.3234	Train Acc: 0.9147	Val Acc: 0.9119
Epoch 43/100	Train Loss: 0.3034	Val Loss: 0.3225	Train Acc: 0.9150	Val Acc: 0.9115
Epoch 44/100	Train Loss: 0.3024	Val Loss: 0.3215	Train Acc: 0.9152	Val Acc: 0.9123
Epoch 45/100	Train Loss: 0.3016	Val Loss: 0.3212	Train Acc: 0.9154	Val Acc: 0.9111
Epoch 46/100	Train Loss: 0.3008	Val Loss: 0.3205	Train Acc: 0.9157	Val Acc: 0.9121
Epoch 47/100	Train Loss: 0.3004	Val Loss: 0.3197	Train Acc: 0.9166	Val Acc: 0.9120
Epoch 48/100	Train Loss: 0.2990	Val Loss: 0.3193	Train Acc: 0.9163	Val Acc: 0.9125
Epoch 49/100	Train Loss: 0.2984	Val Loss: 0.3186	Train Acc: 0.9168	Val Acc: 0.9128
Epoch 50/100	Train Loss: 0.2977	Val Loss: 0.3183	Train Acc: 0.9170	Val Acc: 0.9125
Epoch 51/100	Train Loss: 0.2971	Val Loss: 0.3175	Train Acc: 0.9170	Val Acc: 0.9124
Epoch 52/100	Train Loss: 0.2963	Val Loss: 0.3169	Train Acc: 0.9172	Val Acc: 0.9125
Epoch 53/100	Train Loss: 0.2955	Val Loss: 0.3163	Train Acc: 0.9176	Val Acc: 0.9127
Epoch 54/100	Train Loss: 0.2948	Val Loss: 0.3157	Train Acc: 0.9178	Val Acc: 0.9134
Epoch 55/100	Train Loss: 0.2942	Val Loss: 0.3152	Train Acc: 0.9181	Val Acc: 0.9134
Epoch 56/100	Train Loss: 0.2934	Val Loss: 0.3146	Train Acc: 0.9182	Val Acc: 0.9129
Epoch 57/100	Train Loss: 0.2927	Val Loss: 0.3145	Train Acc: 0.9186	Val Acc: 0.9127
Epoch 58/100	Train Loss: 0.2923	Val Loss: 0.3140	Train Acc: 0.9186	Val Acc: 0.9135

```

plt.figure(figsize=(8, 5))
plt.plot(train_losses, label='Training Loss')
plt.plot(val_losses, label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')

```

```
plt.title('Softmax Regression: Training vs Validation Loss')
plt.legend()
plt.show()


plt.figure(figsize=(8, 5))
plt.plot(train_accuracies, label='Training Accuracy')
plt.plot(val_accuracies, label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Softmax Regression: Training vs Validation Accuracy')
plt.legend()
plt.show()

with torch.no_grad():
    logits_test = X_test @ W + b
    y_test_pred = softmax(logits_test)
    y_test_pred_label = y_test_pred.argmax(dim=1)


test_acc = (y_test_pred_label == y_test).float().mean().item()
print(f"\nFinal Test Accuracy: {test_acc:.4f}")

cm = confusion_matrix(y_test, y_test_pred_label)
ConfusionMatrixDisplay(cm).plot(cmap='Blues')
plt.title("Softmax Regression - Confusion Matrix (Test Set)")
plt.show()
```

Softmax Regression: Training vs Validation Loss




Softmax Regression: Training vs Validation Accuracy



Final Test Accuracy: 0.9190

Softmax Regression - Confusion Matrix (Test Set)



```

classes = list(range(10))
per_class_acc = []
for c in classes:
    mask = (y_test == c)
    acc_c = (y_test_pred_label[mask] == y_test[mask]).float().mean().item()
    per_class_acc.append(acc_c)
    print(f"Accuracy for digit {c}: {acc_c:.4f}")


plt.figure(figsize=(8, 5))
plt.bar(classes, per_class_acc, color='skyblue')
plt.xlabel('Digit')
plt.ylabel('Accuracy')
plt.title('Per-Class Accuracy (Manual Softmax Regression)')
plt.ylim(0.8, 1.0)
plt.show()

```

```

Accuracy for digit 0: 0.9620
Accuracy for digit 1: 0.9741
Accuracy for digit 2: 0.8950
Accuracy for digit 3: 0.8825
Accuracy for digit 4: 0.9289
Accuracy for digit 5: 0.8653
Accuracy for digit 6: 0.9603
Accuracy for digit 7: 0.9314
Accuracy for digit 8: 0.8949
Accuracy for digit 9: 0.8840

```



```

model_torch = nn.Sequential(
    nn.Linear(784, 10)
)

criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model_torch.parameters(), lr=0.01)

epochs_builtin = 100
for epoch in range(epochs_builtin):
    total_loss = 0
    correct = 0
    total = 0
    for X_batch, y_batch in train_loader:
        outputs = model_torch(X_batch)
        loss = criterion(outputs, y_batch)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        total_loss += loss.item()
        preds = outputs.argmax(dim=1)
        correct += (preds == y_batch).sum().item()
        total += y_batch.size(0)

```

```

print(f"[PyTorch Model] Epoch {epoch+1}/{epochs_builtin} | "
      f"Loss: {total_loss/len(train_loader):.4f} | Train Acc: {correct/total:.4f}")

with torch.no_grad():
    outputs_test = model_torch(X_test)
    preds_test = outputs_test.argmax(dim=1)
    acc_torch_test = (preds_test == y_test).float().mean().item()

print(
    f"\nPyTorch Built-in Softmax Regression - Test Accuracy: {acc_torch_test:.4f}")

```

[PyTorch Model] Epoch	1/100	Loss:	1.1765	Train Acc:	0.7682
[PyTorch Model]	Epoch 2/100	Loss:	0.6594	Train Acc:	0.8532
[PyTorch Model]	Epoch 3/100	Loss:	0.5473	Train Acc:	0.8674
[PyTorch Model]	Epoch 4/100	Loss:	0.4933	Train Acc:	0.8761
[PyTorch Model]	Epoch 5/100	Loss:	0.4601	Train Acc:	0.8811
[PyTorch Model]	Epoch 6/100	Loss:	0.4375	Train Acc:	0.8846
[PyTorch Model]	Epoch 7/100	Loss:	0.4203	Train Acc:	0.8876
[PyTorch Model]	Epoch 8/100	Loss:	0.4068	Train Acc:	0.8903
[PyTorch Model]	Epoch 9/100	Loss:	0.3961	Train Acc:	0.8926
[PyTorch Model]	Epoch 10/100	Loss:	0.3872	Train Acc:	0.8943
[PyTorch Model]	Epoch 11/100	Loss:	0.3795	Train Acc:	0.8958
[PyTorch Model]	Epoch 12/100	Loss:	0.3727	Train Acc:	0.8972
[PyTorch Model]	Epoch 13/100	Loss:	0.3669	Train Acc:	0.8983
[PyTorch Model]	Epoch 14/100	Loss:	0.3619	Train Acc:	0.8998
[PyTorch Model]	Epoch 15/100	Loss:	0.3571	Train Acc:	0.9005
[PyTorch Model]	Epoch 16/100	Loss:	0.3533	Train Acc:	0.9016
[PyTorch Model]	Epoch 17/100	Loss:	0.3491	Train Acc:	0.9027
[PyTorch Model]	Epoch 18/100	Loss:	0.3459	Train Acc:	0.9037
[PyTorch Model]	Epoch 19/100	Loss:	0.3426	Train Acc:	0.9042
[PyTorch Model]	Epoch 20/100	Loss:	0.3395	Train Acc:	0.9051
[PyTorch Model]	Epoch 21/100	Loss:	0.3369	Train Acc:	0.9059
[PyTorch Model]	Epoch 22/100	Loss:	0.3343	Train Acc:	0.9061
[PyTorch Model]	Epoch 23/100	Loss:	0.3321	Train Acc:	0.9072
[PyTorch Model]	Epoch 24/100	Loss:	0.3296	Train Acc:	0.9080
[PyTorch Model]	Epoch 25/100	Loss:	0.3279	Train Acc:	0.9079
[PyTorch Model]	Epoch 26/100	Loss:	0.3257	Train Acc:	0.9084
[PyTorch Model]	Epoch 27/100	Loss:	0.3240	Train Acc:	0.9090
[PyTorch Model]	Epoch 28/100	Loss:	0.3225	Train Acc:	0.9097
[PyTorch Model]	Epoch 29/100	Loss:	0.3204	Train Acc:	0.9106
[PyTorch Model]	Epoch 30/100	Loss:	0.3191	Train Acc:	0.9106
[PyTorch Model]	Epoch 31/100	Loss:	0.3173	Train Acc:	0.9113
[PyTorch Model]	Epoch 32/100	Loss:	0.3161	Train Acc:	0.9111
[PyTorch Model]	Epoch 33/100	Loss:	0.3148	Train Acc:	0.9118
[PyTorch Model]	Epoch 34/100	Loss:	0.3131	Train Acc:	0.9124
[PyTorch Model]	Epoch 35/100	Loss:	0.3122	Train Acc:	0.9123
[PyTorch Model]	Epoch 36/100	Loss:	0.3108	Train Acc:	0.9129
[PyTorch Model]	Epoch 37/100	Loss:	0.3098	Train Acc:	0.9129
[PyTorch Model]	Epoch 38/100	Loss:	0.3087	Train Acc:	0.9132
[PyTorch Model]	Epoch 39/100	Loss:	0.3077	Train Acc:	0.9137
[PyTorch Model]	Epoch 40/100	Loss:	0.3064	Train Acc:	0.9142
[PyTorch Model]	Epoch 41/100	Loss:	0.3055	Train Acc:	0.9139
[PyTorch Model]	Epoch 42/100	Loss:	0.3044	Train Acc:	0.9149
[PyTorch Model]	Epoch 43/100	Loss:	0.3035	Train Acc:	0.9148
[PyTorch Model]	Epoch 44/100	Loss:	0.3027	Train Acc:	0.9148
[PyTorch Model]	Epoch 45/100	Loss:	0.3017	Train Acc:	0.9156
[PyTorch Model]	Epoch 46/100	Loss:	0.3008	Train Acc:	0.9160
[PyTorch Model]	Epoch 47/100	Loss:	0.3001	Train Acc:	0.9158
[PyTorch Model]	Epoch 48/100	Loss:	0.2994	Train Acc:	0.9161
[PyTorch Model]	Epoch 49/100	Loss:	0.2985	Train Acc:	0.9166
[PyTorch Model]	Epoch 50/100	Loss:	0.2976	Train Acc:	0.9170
[PyTorch Model]	Epoch 51/100	Loss:	0.2968	Train Acc:	0.9169
[PyTorch Model]	Epoch 52/100	Loss:	0.2962	Train Acc:	0.9174
[PyTorch Model]	Epoch 53/100	Loss:	0.2956	Train Acc:	0.9173
[PyTorch Model]	Epoch 54/100	Loss:	0.2948	Train Acc:	0.9176
[PyTorch Model]	Epoch 55/100	Loss:	0.2942	Train Acc:	0.9181
[PyTorch Model]	Epoch 56/100	Loss:	0.2935	Train Acc:	0.9180
[PyTorch Model]	Epoch 57/100	Loss:	0.2930	Train Acc:	0.9187
[PyTorch Model]	Epoch 58/100	Loss:	0.2922	Train Acc:	0.9186

```

import matplotlib.pyplot as plt
classes = list(range(10))
per_class_acc_torch = []

for c in classes:
    mask = (y_test == c)
    acc_c = (preds_test[mask] == y_test[mask]).float().mean().item()
    per_class_acc_torch.append(acc_c)
    print(f"Digit {c}: {acc_c:.4f}")

# Optional: visualize as bar chart

```

```


plt.figure(figsize=(8, 5))
plt.bar(classes, per_class_acc_torch, color='orange')
plt.xlabel('Digit')
plt.ylabel('Accuracy')
plt.title('Per-Class Accuracy (PyTorch Built-in Softmax Regression)')
plt.ylim(0.8, 1.0)
plt.show()

```

```

Digit 0: 0.9612
Digit 1: 0.9733
Digit 2: 0.8925
Digit 3: 0.8817
Digit 4: 0.9289
Digit 5: 0.8635
Digit 6: 0.9595
Digit 7: 0.9274
Digit 8: 0.8974
Digit 9: 0.8857

```



```

manual_acc = np.array(per_class_acc)
torch_acc = np.array(per_class_acc_torch)

# --- Plot comparison ---
plt.figure(figsize=(10, 6))
bar_width = 0.35
x = np.arange(len(classes))


plt.bar(x - bar_width/2, manual_acc, width=bar_width,
        label='Manual Softmax', color='skyblue')
plt.bar(x + bar_width/2, torch_acc, width=bar_width,
        label='PyTorch Built-in', color='orange')

plt.xlabel('Digit Class')
plt.ylabel('Accuracy')
plt.title('Comparison of Per-Class Accuracy: Manual vs PyTorch Softmax Regression')
plt.xticks(x, classes)
plt.ylim(0.8, 1.0)
plt.legend()
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.show()

# --- Print numerical summary ---
print("Digit | Manual Acc | PyTorch Acc | Difference")
print("-----")
for c in range(10):
    diff = abs(manual_acc[c] - torch_acc[c])
    print(f" {c:>2} | {manual_acc[c]:.4f} | {torch_acc[c]:.4f} | {diff:.4f}")

best_softmax_accuracy = max(val_accuracies)

```



▼ neural network

```

class ManualFeedForward(nn.Module):
    def __init__(self, layers):
        super().__init__()
        self.layers = nn.ModuleList()
        for i in range(len(layers) - 1):
            l = nn.Linear(layers[i], layers[i + 1])
            self.layers.append(l)

        for i, l in enumerate(self.layers):
            if i < len(self.layers) - 1:
                init.kaiming_uniform_(l.weight, nonlinearity='relu')
            else:
                init.xavier_uniform_(l.weight)
                init.zeros_(l.bias)

    def forward(self, x):
        if x.dim() > 2:
            x = x.view(x.size(0), -1)
        self.activations = [x]
        self.z_values = []
        out = x
        for i, layer in enumerate(self.layers):
            z = out @ layer.weight.T + layer.bias
            self.z_values.append(z)
            if i < len(self.layers) - 1:
                out = torch.relu(z)
            else:
                out = z
        self.activations.append(out)
        return out

```

```

def manual_backward(self, y_true, lr):
    batch_size = y_true.size(0)
    logits = self.activations[-1]

    exp_logits = torch.exp(logits - logits.max(dim=1, keepdim=True)[0])
    probs = exp_logits / exp_logits.sum(dim=1, keepdim=True)
    # ensure same device as logits
    y_onehot = torch.zeros_like(probs, device=probs.device)
    y_onehot[torch.arange(batch_size, device=probs.device), y_true] = 1.0
    loss = -torch.sum(y_onehot * torch.log(probs + 1e-9)) / batch_size

    delta = (probs - y_onehot) / batch_size

    for i in reversed(range(len(self.layers))):
        a_prev = self.activations[i]
        z = self.z_values[i]
        dW = delta.T @ a_prev
        db = delta.sum(dim=0)
        self.layers[i].weight.grad = dW
        self.layers[i].bias.grad = db

        if i > 0:
            delta = delta @ self.layers[i].weight
            z_prev = self.z_values[i - 1]
            delta = delta * (z_prev > 0)

    with torch.no_grad():
        for layer in self.layers:
            layer.weight -= lr * layer.weight.grad
            layer.bias -= lr * layer.bias.grad

    return loss.item()

```

```

def train_manual(model, dataloader, lr, device):
    model.train()
    total_loss, total_correct, N = 0, 0, 0
    for X, y in dataloader:
        X, y = X.to(device), y.to(device)
        logits = model.forward(X)
        loss = model.manual_backward(y, lr)
        preds = logits.argmax(dim=1)
        total_correct += (preds == y).sum().item()
        total_loss += loss * X.size(0)
        N += X.size(0)
    return total_loss / N, total_correct / N

```

```

def validate_manual(model, dataloader, device):
    model.eval()
    total_loss, total_correct, N = 0, 0, 0
    with torch.no_grad():
        for X, y in dataloader:
            X, y = X.to(device), y.to(device)
            logits = model.forward(X)
            exp_logits = torch.exp(logits - logits.max(dim=1, keepdim=True)[0])
            probs = exp_logits / exp_logits.sum(dim=1, keepdim=True)
            loss = -torch.log(probs[torch.arange(y.size(0)), y] + 1e-9).mean()
            preds = logits.argmax(dim=1)
            total_correct += (preds == y).sum().item()
            total_loss += loss.item() * X.size(0)
            N += X.size(0)
    return total_loss / N, total_correct / N

```

```

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
INPUT_SIZE = 784
HIDDEN1 = 512
HIDDEN2 = 128
OUTPUT = 10

learning_rates = [0.001, 0.01, 0.1, 1.0]
BATCH_SIZE = 64
PATIENCE = 3
CHECKPOINT_PATH = "best_manual_model.pth"
SEED = 42

```

```

transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])
full_train_dataset = datasets.MNIST(
    './data', train=True, download=True, transform=transform)
train_ds, val_ds = random_split(full_train_dataset, [50000, 10000])
train_loader = DataLoader(train_ds, batch_size=BATCH_SIZE, shuffle=True)
val_loader = DataLoader(val_ds, batch_size=BATCH_SIZE, shuffle=False)

test_dataset = TensorDataset(X_test, y_test)
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False)

model = ManualFeedForward([INPUT_SIZE, HIDDEN1, HIDDEN2, OUTPUT]).to(device)
train_losses, val_losses, train_accs, val_accs = [], [], [], []
best_val_loss = float('inf')
epochs_no_improve = 0
best_model_state = None
epoch = 0
best_epoch = 0

```

```

def train_manual_model(model, train_loader, val_loader, lr, patience=3, device='cpu'):
    model.to(device)
    train_losses, val_losses, train_accs, val_accs = [], [], [], []
    best_val_loss = float('inf')
    epochs_no_improve = 0
    best_model_state = None
    epoch = 0
    best_epoch = 0

    # get run time
    start_time = time.time()
    while epoch < 40:
        epoch += 1
        train_loss, train_acc = train_manual(model, train_loader, lr, device)
        val_loss, val_acc = validate_manual(model, val_loader, device)

        train_losses.append(train_loss)
        val_losses.append(val_loss)
        train_accs.append(train_acc)
        val_accs.append(val_acc)

        print(f"Epoch {epoch:02d} | Train Loss: {train_loss:.4f} | Train Acc: {train_acc*100:.2f}% | "
              f"Val Loss: {val_loss:.4f} | Val Acc: {val_acc*100:.2f}%")

        # Early stopping
        if val_loss < best_val_loss:
            best_val_loss = val_loss
            best_model_state = {k: v.clone()
                                for k, v in model.state_dict().items()}
            best_epoch = epoch
            epochs_no_improve = 0
        else:
            epochs_no_improve += 1
            if epochs_no_improve >= patience:
                print(f"Early stopping triggered after {epoch} epochs.")
                break

    end_time = time.time()
    run_time = end_time - start_time
    # Restore best model
    if best_model_state is not None:
        model.load_state_dict(best_model_state)

    return {
        "train_losses": train_losses,
        "val_losses": val_losses,
        "train_accs": train_accs,
        "val_accs": val_accs,
        "best_val_loss": best_val_loss,
        "best_model_state": best_model_state,
        "best_epoch": best_epoch,
        "best_run_time": run_time
    }

```

```

learning_rates = [0.001, 0.01, 0.1, 1.0]

lr_results = {}

for LR in learning_rates:
    print(f"\n--- Training with LR={LR} ---")
    model = ManualFeedForward([INPUT_SIZE, HIDDEN1, HIDDEN2, OUTPUT])
    result = train_manual_model(
        model, train_loader, val_loader, lr=LR, patience=PATIENCE, device=device)
    lr_results[LR] = result

# find best learning rate
best_lr = 0.01
print(best_lr)
best_lr_val_error = lr_results[best_lr]["best_val_loss"]
print(
    f"\nBest Learning Rate: {best_lr} with Val Loss: {lr_results[best_lr]['best_val_loss']:.4f}")
print(
    f"Model with batch size {BATCH_SIZE}, number of layers 3, number of neurons [{HIDDEN1}, {HIDDEN2}]")

# Plot learning curves
plt.figure(figsize=(10, 5))
for LR in learning_rates:
    plt.plot(lr_results[LR]["train_losses"], label=f'Train Loss LR={LR}')
    plt.plot(lr_results[LR]["val_losses"], '--', label=f'Val Loss LR={LR}')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Learning Rate Analysis')
plt.legend()
plt.show()

best_nn_accuracy = max(lr_results[best_lr]["val_accs"])
best_nn_run_time = lr_results[best_lr]["best_run_time"]

```



```

--- Training with LR=0.001 ---
Epoch 01 | Train Loss: 1.1035 | Train Acc: 67.30% | Val Loss: 0.6466 | Val Acc: 81.61%
Epoch 02 | Train Loss: 0.5327 | Train Acc: 85.08% | Val Loss: 0.4666 | Val Acc: 86.46%
Epoch 03 | Train Loss: 0.4167 | Train Acc: 88.19% | Val Loss: 0.3943 | Val Acc: 88.40%
Epoch 04 | Train Loss: 0.3606 | Train Acc: 89.73% | Val Loss: 0.3531 | Val Acc: 89.66%
Epoch 05 | Train Loss: 0.3256 | Train Acc: 90.68% | Val Loss: 0.3250 | Val Acc: 90.44%
Epoch 06 | Train Loss: 0.3008 | Train Acc: 91.36% | Val Loss: 0.3051 | Val Acc: 90.91%
Epoch 07 | Train Loss: 0.2816 | Train Acc: 91.91% | Val Loss: 0.2888 | Val Acc: 91.46%
Epoch 08 | Train Loss: 0.2662 | Train Acc: 92.33% | Val Loss: 0.2752 | Val Acc: 92.02%
Epoch 09 | Train Loss: 0.2533 | Train Acc: 92.73% | Val Loss: 0.2647 | Val Acc: 92.21%
Epoch 10 | Train Loss: 0.2422 | Train Acc: 93.04% | Val Loss: 0.2544 | Val Acc: 92.70%
Epoch 11 | Train Loss: 0.2327 | Train Acc: 93.36% | Val Loss: 0.2467 | Val Acc: 92.83%
Epoch 12 | Train Loss: 0.2241 | Train Acc: 93.65% | Val Loss: 0.2387 | Val Acc: 93.17%
Epoch 13 | Train Loss: 0.2164 | Train Acc: 93.86% | Val Loss: 0.2325 | Val Acc: 93.41%
Epoch 14 | Train Loss: 0.2093 | Train Acc: 94.04% | Val Loss: 0.2263 | Val Acc: 93.58%
Epoch 15 | Train Loss: 0.2029 | Train Acc: 94.21% | Val Loss: 0.2207 | Val Acc: 93.82%
Epoch 16 | Train Loss: 0.1970 | Train Acc: 94.38% | Val Loss: 0.2162 | Val Acc: 93.89%
Epoch 17 | Train Loss: 0.1915 | Train Acc: 94.56% | Val Loss: 0.2108 | Val Acc: 94.06%
Epoch 18 | Train Loss: 0.1864 | Train Acc: 94.73% | Val Loss: 0.2065 | Val Acc: 94.21%
Epoch 19 | Train Loss: 0.1817 | Train Acc: 94.84% | Val Loss: 0.2024 | Val Acc: 94.31%
Epoch 20 | Train Loss: 0.1771 | Train Acc: 95.01% | Val Loss: 0.1987 | Val Acc: 94.39%
Epoch 21 | Train Loss: 0.1730 | Train Acc: 95.10% | Val Loss: 0.1963 | Val Acc: 94.38%
Epoch 22 | Train Loss: 0.1689 | Train Acc: 95.25% | Val Loss: 0.1919 | Val Acc: 94.40%
Epoch 23 | Train Loss: 0.1651 | Train Acc: 95.36% | Val Loss: 0.1891 | Val Acc: 94.57%
Epoch 24 | Train Loss: 0.1616 | Train Acc: 95.47% | Val Loss: 0.1858 | Val Acc: 94.73%
Epoch 25 | Train Loss: 0.1582 | Train Acc: 95.56% | Val Loss: 0.1823 | Val Acc: 94.84%
Epoch 26 | Train Loss: 0.1549 | Train Acc: 95.62% | Val Loss: 0.1797 | Val Acc: 94.92%
Epoch 27 | Train Loss: 0.1517 | Train Acc: 95.74% | Val Loss: 0.1770 | Val Acc: 94.97%
Epoch 28 | Train Loss: 0.1487 | Train Acc: 95.84% | Val Loss: 0.1753 | Val Acc: 94.96%
Epoch 29 | Train Loss: 0.1459 | Train Acc: 95.92% | Val Loss: 0.1719 | Val Acc: 95.14%
Epoch 30 | Train Loss: 0.1431 | Train Acc: 95.98% | Val Loss: 0.1700 | Val Acc: 95.17%
Epoch 31 | Train Loss: 0.1405 | Train Acc: 96.07% | Val Loss: 0.1674 | Val Acc: 95.36%
Epoch 32 | Train Loss: 0.1379 | Train Acc: 96.12% | Val Loss: 0.1654 | Val Acc: 95.39%
Epoch 33 | Train Loss: 0.1355 | Train Acc: 96.22% | Val Loss: 0.1635 | Val Acc: 95.45%
Epoch 34 | Train Loss: 0.1331 | Train Acc: 96.29% | Val Loss: 0.1613 | Val Acc: 95.51%
Epoch 35 | Train Loss: 0.1308 | Train Acc: 96.39% | Val Loss: 0.1599 | Val Acc: 95.57%
Epoch 36 | Train Loss: 0.1286 | Train Acc: 96.46% | Val Loss: 0.1582 | Val Acc: 95.52%
Epoch 37 | Train Loss: 0.1264 | Train Acc: 96.53% | Val Loss: 0.1563 | Val Acc: 95.59%
Epoch 38 | Train Loss: 0.1243 | Train Acc: 96.59% | Val Loss: 0.1542 | Val Acc: 95.67%
Epoch 39 | Train Loss: 0.1223 | Train Acc: 96.63% | Val Loss: 0.1529 | Val Acc: 95.64%
Epoch 40 | Train Loss: 0.1203 | Train Acc: 96.70% | Val Loss: 0.1514 | Val Acc: 95.73%

```

```

--- Training with LR=0.01 ---
Epoch 01 | Train Loss: 0.4241 | Train Acc: 87.53% | Val Loss: 0.2612 | Val Acc: 92.44%
Epoch 02 | Train Loss: 0.2110 | Train Acc: 93.81% | Val Loss: 0.2044 | Val Acc: 94.20%
Epoch 03 | Train Loss: 0.1652 | Train Acc: 95.25% | Val Loss: 0.1764 | Val Acc: 94.80%
Epoch 04 | Train Loss: 0.1383 | Train Acc: 96.07% | Val Loss: 0.1601 | Val Acc: 95.35%
Epoch 05 | Train Loss: 0.1185 | Train Acc: 96.58% | Val Loss: 0.1494 | Val Acc: 95.72%
Epoch 06 | Train Loss: 0.1037 | Train Acc: 97.10% | Val Loss: 0.1353 | Val Acc: 95.94%
Epoch 07 | Train Loss: 0.0919 | Train Acc: 97.48% | Val Loss: 0.1279 | Val Acc: 96.29%

```

```

print("\n--- Evaluating best model (LR = 0.01) on test set ---")

best_model = ManualFeedForward([INPUT_SIZE, HIDDEN1, HIDDEN2, OUTPUT])

best_model.load_state_dict(lr_results[0.01]["best_model_state"])
best_model.to(device)
best_model.eval()

total_loss, total_correct, total_samples = 0, 0, 0

with torch.no_grad():
    for X_batch, y_batch in test_loader:
        X_batch, y_batch = X_batch.to(device), y_batch.to(device)
        outputs = best_model(X_batch)

        exp_logits = torch.exp(outputs - outputs.max(dim=1, keepdim=True)[0])
        probs = exp_logits / exp_logits.sum(dim=1, keepdim=True)
        loss = -torch.log(probs[torch.arange(y_batch.size(0)), y_batch] + 1e-9).mean()

        total_loss += loss.item() * X_batch.size(0)
        preds = outputs.argmax(dim=1)
        total_correct += (preds == y_batch).sum().item()
        total_samples += X_batch.size(0)

avg_test_loss = total_loss / total_samples
test_accuracy = 100 * total_correct / total_samples

print(f"Test Loss: {avg_test_loss:.4f}")
print(f"Test Accuracy: {test_accuracy:.2f}%")

```

```
-- Training with LR=0.1 --
Epoch 01 | Train Loss: 0.2264 | Train Acc: 93.05% | Val Loss: 0.2008 | Val Acc: 94.07%
Epoch 02 | Train Loss: 0.0652 | Train Acc: 97.49% | Val Loss: 0.0906 | Val Acc: 97.15%
Epoch 03 | Train Loss: 0.0497 | Train Acc: 98.52% | Val Loss: 0.0839 | Val Acc: 97.39%
Epoch 04 | Train Loss: 0.0306 | Train Acc: 99.15% | Val Loss: 0.0896 | Val Acc: 97.30%
Epoch 05 | Train Loss: 0.0187 | Train Acc: 99.52% | Val Loss: 0.1119 | Val Acc: 96.71%
```

```
batch_sizes = [16, 32, 64, 128]
batch_results = {}

for BS in batch_sizes:
    print(f"\n--- Training with batch size={BS} ---")
    train_loader_bs = DataLoader(train_ds, batch_size=BS, shuffle=True)
    val_loader_bs = DataLoader(val_ds, batch_size=BS, shuffle=False)

    model = ManualFeedForward([INPUT_SIZE, HIDDEN1, HIDDEN2, OUTPUT])
    result = train_manual_model(
        model, train_loader_bs, val_loader_bs, lr=0.01, patience=PATIENCE, device=device)
    batch_results[BS] = result

# Find the best batch size based on validation loss
best_bs = min(batch_results, key=lambda bs: batch_results[bs]["best_val_loss"])
best_bs_val_error = batch_results[best_bs]["best_val_loss"]
print(
    f"\nBest Batch Size: {best_bs} with Val Loss: {batch_results[best_bs]['best_val_loss']:.4f}")
print(
    f"meet with learning rate 0.01, number of layers 3, number of neurons [{HIDDEN1}, {HIDDEN2}]")
```

```
# Plot batch size comparison
plt.figure(figsize=(10, 5))
for BS in batch_sizes:
    plt.plot(batch_results[BS]["train_losses"], label=f'Train Loss BS={BS}')
    plt.plot(batch_results[BS]["val_losses"], '--', label=f'Val Loss BS={BS}')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Batch Size Analysis')
plt.legend()
plt.show()
```

